

SOC2025 – Assignment 1

Handwritten Digit Recognition using CNN

Submitted by
Esha Yindukuri
24B2119

Objective

Building a model which can recognize handwritten digits from the **MNIST** dataset using a **Convolutional Neural Network (CNN)**, involving layers such as **convolutional, pooling, and fully connected layers** to automatically extract features from the images and classify the digits (0-9)

Dataset

We used the **MNIST dataset**, which consists of:

- **60,000 training images** and **10,000 test images**
- Each image is a **28×28 grayscale** image
- Labels represent digits from **0 to 9**
- Each pixel has a value between 0 and 255, where 0 is black and 255 is white.

The dataset was directly loaded from **TensorFlow's Keras API**.

Data Preprocessing

Before training the model, cleaned and prepared the data using these steps:

1. **Normalization:**
Scaled the pixel values from 0–255 down to **0–1** by dividing each pixel by 255. This helps the model learn faster.
2. **Reshaping:**
Reshaped the images to have one more dimension (28×28×1) so that they can be used in a CNN model.
3. **One-Hot Encoding:**
The labels (like 0, 1, 2, ..., 9) were changed into a format suitable for classification using one-hot vectors. For example, the digit “3” becomes [0, 0, 0, 1, 0, 0, 0, 0, 0].

CNN Model Architecture

We created the CNN model using the **Sequential API** in TensorFlow Keras. The model includes several layers:

1. **First Convolutional Layer**
 - 32 filters of size 3×3
 - Uses ReLU activation
 - Finds simple shapes and edges
2. **First MaxPooling Layer**
 - Pool size 2×2
 - Reduces the size of the feature map
3. **Second Convolutional Layer**
 - 64 filters of size 3×3
 - ReLU activation again
 - Detects more complex patterns
4. **Second MaxPooling Layer**
 - Pool size 2×2
 - Further reduces size
5. **Flatten Layer**
 - Converts 2D output into a 1D vector
6. **Fully Connected Dense Layer**
 - 128 neurons
 - ReLU activation
 - Learns important combinations of features
7. **Dropout Layer**
 - 50% of neurons turned off randomly during training to prevent overfitting
8. **Output Layer**
 - 10 neurons (for digits 0–9)
 - Softmax activation to give probability for each class

Model Compilation

Before training, we compiled the model with these settings:

- **Optimizer:** Adam (a popular optimizer that adjusts learning rate automatically)
- **Loss Function:** Categorical Crossentropy (used for multi-class classification)
- **Metric:** Accuracy (to check how many predictions are correct)

Training the Model

Trained the model using:

- **10 epochs** (10 times through the full training data)
- **Batch size of 128** (128 images are processed before updating the weights)
- **Validation split of 20%** (part of training data used to check performance after every epoch)

During training, tracked both **training and validation accuracy/loss**, and plotted graphs to check how the model was learning.

Result

After training, tested the model on the test set and got an accuracy of around **99.20%**. Also

- Made predictions on a few test images and showed both the real and predicted labels.
- Plotted a **confusion matrix** to see which digits were getting confused with each other.

